

YADLAPALLI SANKAR SAI NARAYANA

sankarsainarayana@gmail.com | +91-7207404856 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

Results-driven AI and Cloud Software Engineer with 1+ year of hands-on experience delivering production-grade LLM applications, AI agents, and Retrieval-Augmented Generation (RAG) systems. Certified Kubernetes Administrator (CKA) specializing in cloud-native architectures on AWS — including containerization, orchestration, CI/CD pipelines, and high-availability deployments. Built and deployed an intelligent MCP-based observability platform integrating 11 tools that reduced mean time to resolution (MTTR) by ~40%. Adept at bridging AI engineering with DevOps practices to deliver reliable, scalable, and cost-efficient solutions in fast-paced production environments.

TECHNICAL SKILLS

Programming Languages: Python, TypeScript

AI/ML & LLM Frameworks: LangChain, LangGraph, TensorFlow, Keras, PyTorch, Scikit-learn, HuggingFace

AI Tools & Agentic Development: Claude Code, OpenAI Codex, Prompt Engineering, RAG, Vector Databases, MCP (Model Context Protocol)

Cloud & DevOps: AWS (EKS, EC2, IAM, S3), Docker, Kubernetes, Helm, GitHub Actions, CI/CD, Linux

Observability & Monitoring: Distributed Tracing, Log Analysis, Incident Management, Root Cause Analysis (RCA), On-Call Support

Backend & APIs: FastAPI, Flask, FastMCP, REST APIs

Databases: MySQL, PostgreSQL

BI & Visualization: Microsoft Power BI, ETL Pipelines, Jupyter Notebook

PROFESSIONAL EXPERIENCE

DevOps Engineer Intern | Yellow AI (Yellow Messenger)

March 2025 – Present

Bangalore, India

Cloud Infrastructure & Site Reliability

- **Managed and maintained** production Kubernetes clusters on AWS EKS across **3+ environments**, performing version upgrades, patching, and compatibility checks that reduced deployment failure rate by **~30%**.
- **Led on-call incident response** for production systems, resolving **15+ critical incidents** and reducing mean time to resolution (MTTR) by **~40%** through structured log analysis and rapid remediation.
- **Optimized resource utilization** across containerized workloads, achieving **~25% reduction** in infrastructure overhead by identifying and resolving CPU/memory bottlenecks in Kubernetes deployments.
- **Accelerated release cycles by ~35%** by building and managing scalable CI/CD pipelines using GitHub Actions, Docker, and Kubernetes (EKS), enabling faster and more reliable deployments.
- **Enforced security and compliance** through IAM policy management, cluster RBAC configuration, and routine patching across AWS infrastructure components.

AI Platform & Observability Engineering

- **Architected and deployed** an intelligent on-demand tool-calling system for observability and incident analysis, integrating **11 MCP tools** into a single unified diagnostic platform used by the engineering team daily.
- **Engineered an LLM-orchestrated automation layer** that interprets natural language queries (e.g., trace ID, region) and autonomously selects and executes the correct observability tool, eliminating **~3 hours/day** of manual debugging effort.
- **Implemented autonomous root cause analysis (RCA)** using LangChain-based orchestration, enabling the system to diagnose distributed trace failures and surface actionable recommendations without human intervention.
- **Reduced mean time to detect (MTTD) by ~45%** by transforming complex multi-step observability workflows into a single-query automated diagnostic system.
- **Improved developer productivity by ~50%** by eliminating repetitive log-diving and tool-switching, consolidating insights from distributed traces, failures, and anomalies into one interface.

KEY PROJECTS

Intelligent Observability Platform (MCP + LLM) | *Yellow AI*

- Designed a production-grade agentic system using FastMCP and LangChain that routes observability queries to 11 specialized tools, cutting incident investigation time from ~60 minutes to under 5 minutes.
- Implemented structured decision-making logic with LLM-based orchestration (Claude/OpenAI), enabling zero-touch automated RCA for distributed microservice failures.

RAG-Based Knowledge Assistant

- Built a Retrieval-Augmented Generation (RAG) pipeline using LangChain, vector databases, and HuggingFace embeddings, achieving ~85% answer relevance on domain-specific Q&A benchmarks.
- Deployed the solution with FastAPI and Docker, enabling sub-second retrieval latency across a corpus of 10,000+ documents.

EDUCATION

B.Tech in Information Technology (Honors)

November 2021 – May 2025

JNTU-GV College of Engineering, Vizianagaram, Andhra Pradesh | GPA: 8.0+/10

CERTIFICATIONS

Certified Kubernetes Administrator (CKA) — Linux Foundation, USA [\[View Certificate\]](#)